

▼ Importing libraries and datasets

```
import numpy as np
import pandas as pd
import datetime as dt
import numpy as np
import matplotlib.pyplot as plt
import xgboost as xgb
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score,
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
from imblearn.over_sampling import SMOTE
from xgboost import plot_importance
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import KFold # import KFold
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn import metrics
```

```
df_response = pd.read_csv('Retail_Data_Response.csv')
df_transactions = pd.read_csv('Retail_Data_Transactions.csv', parse_dates=['trans_c
```

```
df_response.head()
```

	customer_id	response
0	CS1112	0
1	CS1113	0
2	CS1114	1
3	CS1115	1
4	CS1116	1

```
df_transactions.head()
```

	customer_id	trans_date	tran_amount
0	CS5295	2013-02-11	35
1	CS4768	2015-03-15	39
2	CS2122	2013-02-26	52
3	CS1217	2011-11-16	99
4	CS1850	2013-11-20	78

```
print(df_transactions['trans_date'].min())
print(df_transactions['trans_date'].max())
```

```
2011-05-16 00:00:00
2015-03-16 00:00:00
```

▼ Data Preparation

```
## since the last date of the data is 16 March 2015, the campaign date is assumed to be 17 March 2015
## RFM model will be used to predict campaign response. Recency is calculated as the number of days since the last transaction
```

```
campaign_date = dt.datetime(2015,3,17)
df_transactions['recent']= campaign_date - df_transactions['trans_date']
df_transactions['recent'].astype('timedelta64[D]')
df_transactions['recent']=df_transactions['recent'] / np.timedelta64(1, 'D')
df_transactions.head()
```

	customer_id	trans_date	tran_amount	recent
0	CS5295	2013-02-11	35	764.0
1	CS4768	2015-03-15	39	2.0
2	CS2122	2013-02-26	52	749.0
3	CS1217	2011-11-16	99	1217.0
4	CS1850	2013-11-20	78	482.0

```
## create data set with RFM variables
```

```
df_rfm = df_transactions.groupby('customer_id').agg({'recent': lambda x:x.min(),
                                                    'customer_id': lambda x: len(x),
                                                    'tran_amount': lambda x: x.sum()})
```

```
df_rfm.rename(columns={'recent': 'recency',
                      'customer_id': 'frequency',
                      'tran_amount': 'monetary_value'}, inplace=True)
```

```
df_rfm = df_rfm.reset_index()
df_rfm.head()
```


▼ Calculating response rate

```
response_rate = df_response.groupby('response').agg({'customer_id': lambda x: len(x)})
response_rate.head()
```

	response	customer_id
0	0	6237
1	1	647

```
plt.figure(figsize=(5,5))
x=range(2)
plt.bar(x,response_rate['customer_id'])
plt.xticks(response_rate.index)
plt.title('Response Distribution')
plt.xlabel('Convert or Not')
plt.ylabel('no. of users')
plt.show()
```

data is imbalanced



merging two data sets

```
df_modeling = pd.merge(df_response,df_rfm)
df_modeling.head()
```

	customer_id	response	recency	frequency	monetary_value
0	CS1112	0	62.0	15	1012
1	CS1113	0	36.0	20	1490

► Creating train and test dataset

```

# splitting dataframe into X and y

X = df_modeling.drop(columns=['response', 'customer_id'])
y = df_modeling['response']


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

print("Number transactions X_train dataset: ", X_train.shape)
print("Number transactions y_train dataset: ", y_train.shape)
print("Number transactions X_test dataset: ", X_test.shape)
print("Number transactions y_test dataset: ", y_test.shape)

Number transactions X_train dataset: (4818, 3)
Number transactions y_train dataset: (4818,)
Number transactions X_test dataset: (2066, 3)
Number transactions y_test dataset: (2066,)

sns.scatterplot(data=df_modeling, x='recency', y='monetary_value', hue='response')
sns.despine()
plt.title("Imbalanced Data")

```



► Fixing imbalanced with Undersampling

[] ↪ 1 cell hidden

► Fixing imbalanced with Oversampling


[] ↪ 1 cell hidden

▼ Fixing imbalanced with SMOTE

```
sm = SMOTE(random_state=0)
sm.fit(X_train, y_train)
X_SMOTE, y_SMOTE = sm.fit_sample(X_train, y_train)
df_SMOTE = pd.concat([pd.DataFrame(data=X_SMOTE), pd.DataFrame(data=y_SMOTE)], axis=
df_SMOTE.columns= ['recency', 'frequency', 'monetary_value', 'response']
```

```
sns.scatterplot(data=df_SMOTE, x='recency', y='monetary_value', hue='response')
sns.despine()
plt.title("SMOTE Data")
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: warn(msg, category=FutureWarning)
Text(0.5, 1.0, 'SMOTE Data')
```



► Logistic Regression Model

[] ↪ 9 cells hidden

► XGBoost

[] ↪ 10 cells hidden

Improve

▼ Import File and Create Feature

ลองเพิ่ม Feature Age of Transaction แต่พบว่า ไม่ได้ทำให้ AUC เพิ่มขึ้น

```
df_response = pd.read_csv('Retail_Data_Response.csv')
df_transactions = pd.read_csv('Retail_Data_Transactions.csv', parse_dates=['trans_date'])

campaign_date = dt.datetime(2015,3,17)
df_transactions['recent'] = campaign_date - df_transactions['trans_date']
df_transactions['recent'].astype('timedelta64[D]')
df_transactions['recent'] = df_transactions['recent'] / np.timedelta64(1, 'D')
df_transactions.head()
```

#Add More feature

```
campaign_date = dt.datetime(2015,3,17)
df_transactions['first_transaction'] = campaign_date - df_transactions['trans_date']
df_transactions['first_transaction'].astype('timedelta64[D]')
df_transactions['first_transaction'] = df_transactions['first_transaction'] / np.timedelta64(1, 'D')
df_transactions['age_of_T'] = campaign_date - df_transactions['trans_date']
df_transactions['age_of_T'].astype('timedelta64[D]')
df_transactions['age_of_T'] = df_transactions['age_of_T'] / np.timedelta64(1, 'D')
df_transactions.head()
```

	customer_id	trans_date	tran_amount	recent	first_transaction	age_of_T
0	CS5295	2013-02-11	35	764.0	764.0	764.0
1	CS4768	2015-03-15	39	2.0	2.0	2.0
2	CS2122	2013-02-26	52	749.0	749.0	749.0
3	CS1217	2011-11-16	99	1217.0	1217.0	1217.0
4	CS1850	2013-11-20	78	482.0	482.0	482.0

create data set with RFM variables

```
df_rfm = df_transactions.groupby('customer_id').agg({'recent': lambda x: x.min(),
                                                    'customer_id': lambda x: len(x),
                                                    'tran_amount': lambda x: x.sum(),
                                                    'first_transaction': lambda x: x.min(),
                                                    'age_of_T': lambda x: x.max()})

df_rfm.rename(columns={'recent': 'recency',
                      'customer_id': 'frequency',
                      'tran_amount': 'monetary_value',
                      'first_transaction': 'first_transaction',
                      'age_of_T': 'age_of_T'}, inplace=True)
```

```
df_rfm['rfm'] = df_rfm['recency'] * df_rfm['frequency'] * df_rfm['monetary_value']
```

```
df_rfm.head(2)
```

	recency	frequency	monetary_value	first_transaction	age_of_T
customer_id					
CS1112	62.0	15	1012	1371.0	1309.0
CS1113	36.0	20	1490	1390.0	1354.0

```
df_rfm = df_rfm.reset_index()
df_rfm.head()
```

	customer_id	recency	frequency	monetary_value	first_transaction	age_of_
0	CS1112	62.0	15	1012	1371.0	1309.
1	CS1113	36.0	20	1490	1390.0	1354.
2	CS1114	33.0	19	1432	1342.0	1309.
3	CS1115	12.0	22	1659	1315.0	1303.
4	CS1116	204.0	13	857	1359.0	1155.

```
df_modeling = pd.merge(df_response,df_rfm)
df_modeling.head()
```

	customer_id	response	recency	frequency	monetary_value	first_transactio
0	CS1112	0	62.0	15	1012	1371.
1	CS1113	0	36.0	20	1490	1390.
2	CS1114	1	33.0	19	1432	1342.
3	CS1115	1	12.0	22	1659	1315.
4	CS1116	1	204.0	13	857	1359.

```
## splitting dataframe into X and y
```

```
X = df_modeling.drop(columns=['first_transaction','age_of_T','rfm','response','cust
# X = df_modeling[['age_of_T','rfm']]
y = df_modeling['response']
```

```
df_modeling.describe()
```

	response	recency	frequency	monetary_value	first_transaction	
count	6884.000000	6884.000000	6884.000000	6884.000000	6884.000000	6884.000000
mean	0.093986	81.024985	18.153544	1179.892214	1322.488379	1322.488379
std	0.291831	83.251016	5.184476	465.421365	84.553118	84.553118
min	0.000000	1.000000	4.000000	149.000000	507.000000	507.000000

```
df_modeling.columns
```

```
Index(['customer_id', 'response', 'recency', 'frequency', 'monetary_value',
      'first_transaction', 'age_of_T', 'rfm'],
      dtype='object')
```

```
max    1.000000    83.251016    5.184476    465.421365    84.553118    84.553118
```

▼ Explore

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
print("Number transactions X_train dataset: ", X_train.shape)
```

```
print("Number transactions y_train dataset: ", y_train.shape)
```

```
print("Number transactions X_test dataset: ", X_test.shape)
```

```
print("Number transactions y_test dataset: ", y_test.shape)
```

```
Number transactions X_train dataset: (4818, 3)
```

```
Number transactions y_train dataset: (4818,)
```

```
Number transactions X_test dataset: (2066, 3)
```

```
Number transactions y_test dataset: (2066,)
```

```
sm = SMOTE(random_state=0)
```

```
sm.fit(X_train, y_train)
```

```
X_SMOTE, y_SMOTE = sm.fit_sample(X_train, y_train)
```

```
df_SMOTE_imp = pd.concat([pd.DataFrame(data=X_SMOTE),pd.DataFrame(data=y_SMOTE)], axis=1)
```

```
df_SMOTE_imp.columns= ['recency', 'frequency', 'monetary_value', 'response']
```

```
# df_SMOTE_imp.columns= ['age_of_T', 'rfm', 'response']
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning:
  warnings.warn(msg, category=FutureWarning)
```

```
df_SMOTE_imp.head(5)
```


	recency	frequency	monetary_value	response
0	115 0	12 0	564 0	0

```
df_SMOTE_imp['response'].unique()
```

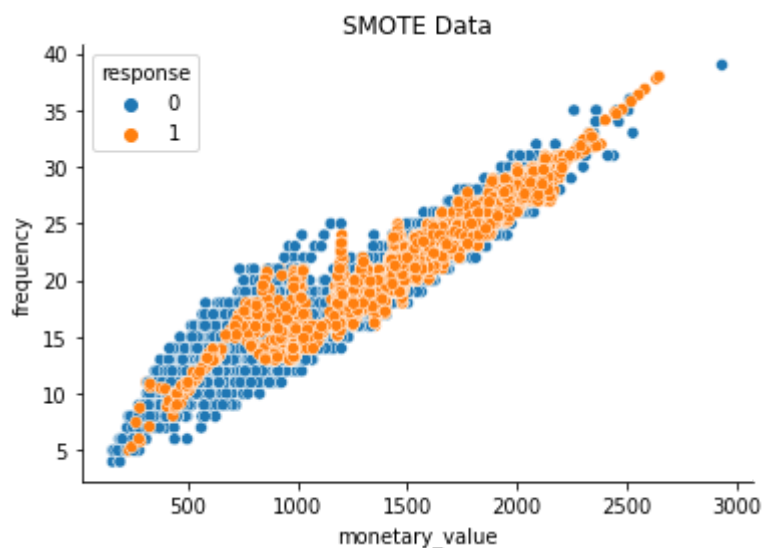
```
array([0, 1])
```

```
df_SMOTE_imp.describe()
```

	recency	frequency	monetary_value	response
count	8778.000000	8778.000000	8778.000000	8778.000000
mean	82.329037	19.546411	1311.604887	0.500000
std	75.408961	4.969402	439.465126	0.500028
min	1.000000	4.000000	157.000000	0.000000
25%	27.395891	16.000000	1048.000000	0.000000
50%	62.762745	20.000000	1376.754041	0.500000
75%	115.295107	23.000000	1607.000000	1.000000
max	679.000000	39.000000	2933.000000	1.000000


```
sns.scatterplot(data=df_SMOTE_imp, x='monetary_value', y='frequency', hue='response')
sns.despine()
plt.title("SMOTE Data")
```

```
Text(0.5, 1.0, 'SMOTE Data')
```




```
sns.scatterplot(data=df_SMOTE_imp, x='recency', y='monetary_value', hue='response')
sns.despine()
plt.title("SMOTE Data")
```

```
Text(0.5, 1.0, 'SMOTE Data')
```



```
sns.scatterplot(data=df_SMOTE_imp, x='recency', y='frequency', hue='response')
sns.despine()
plt.title("SMOTE Data")
```

```
Text(0.5, 1.0, 'SMOTE Data')
```



```
#Check Imbalance
```

```
df_SMOTE_imp.response.value_counts()
```

```
1    4389
0    4389
Name: response, dtype: int64
```

▼ Tuning Model

ลองแบ่ง ข้อมูลด้วย KFold พบว่า n_splits = 12 จะได้ AUC ดีที่สุด

```
#K-Fold Cross validation
kf = KFold(n_splits=12,shuffle=True,random_state=101)
for train_index, test_index in kf.split(X):
    X_train, X_test = X.iloc[train_index,:], X.iloc[test_index,:]
    y_train, y_test = y[train_index], y[test_index]
    print('Train Shape X: {} Y : {}'.format(X_train.shape,y_train.shape))
    print('Test Shape X: {} Y : {}'.format(X_test.shape,y_test.shape))
```

```
#Fix Imbalance with SMOTE
sm = SMOTE(random_state=0)
sm.fit(X_train, y_train)
X_SMOTE, y_SMOTE = sm.fit_sample(X_train, y_train)

#XGb Model
xgb_model = xgb.XGBClassifier(objective='binary:logistic', eval_metric="auc",
    base_score=0.5,
    learning_rate =0.10,
    n_estimators=5000,
    max_depth=5,
    min_child_weight=2,
    gamma=0.1,
    subsample=0.4,
    colsample_bytree=0.4,
    nthread=4)
predicted_y = []
expected_y = []

xgb_model_SMOTE = xgb_model.fit(X_SMOTE, y_SMOTE, early_stopping_rounds=5, eval_s
predictions = xgb_model_SMOTE.predict(X_SMOTE)
predicted_y.extend(predictions)
expected_y.extend(y_SMOTE)
report_train = classification_report(expected_y, predicted_y)
print('training set')
print(report_train)

predicted_y = []
expected_y = []
predictions = xgb_model_SMOTE.predict(X_test.to_numpy())
predicted_y.extend(predictions)
expected_y.extend(y_test)
report_test = classification_report(expected_y, predicted_y)
print('test set')
print(report_test)
```

```
[5]      validation_0-auc:0.76672
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning
warnings.warn(msg, category=FutureWarning)
```

training set

	precision	recall	f1-score	support
0	0.83	0.71	0.77	5715
1	0.75	0.86	0.80	5715
accuracy			0.78	11430
macro avg	0.79	0.78	0.78	11430
weighted avg	0.79	0.78	0.78	11430

test set

	precision	recall	f1-score	support
0	0.95	0.74	0.83	522
1	0.19	0.63	0.29	51

accuracy

0.72

572

accuracy			0.73	573
macro avg	0.57	0.68	0.56	573
weighted avg	0.89	0.73	0.79	573

Train Shape X: (6311, 3) Y : (6311,)

Test Shape X: (573, 3) Y : (573,)

[0] validation_0-auc:0.697561

Will train until validation_0-auc hasn't improved in 5 rounds.

[1] validation_0-auc:0.749126

[2] validation_0-auc:0.739911

[3] validation_0-auc:0.762914

[4] validation_0-auc:0.771904

[5] validation_0-auc:0.783966

[6] validation_0-auc:0.765066

[7] validation_0-auc:0.760806

[8] validation_0-auc:0.766052

[9] validation_0-auc:0.759775

[10] validation_0-auc:0.753991

Stopping. Best iteration:

[5] validation_0-auc:0.783966

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning.warn(msg, category=FutureWarning)

training set

	precision	recall	f1-score	support
0	0.73	0.75	0.74	5706
1	0.74	0.72	0.73	5706
accuracy			0.74	11412
macro avg	0.74	0.74	0.74	11412
weighted avg	0.74	0.74	0.74	11412

test set

	precision	recall	f1-score	support
0	0.97	0.72	0.82	531
1	0.17	0.71	0.27	42

▼ Result

Train Shape X: (6311, 3) Y : (6311,) Test Shape X: (573, 3) Y : (573,)

[5] validation_0-auc:0.783966

training set precision recall f1-score support

0	0.73	0.75	0.74	5706
1	0.74	0.72	0.73	5706

accuracy			0.74	11412
macro avg	0.74	0.74	0.74	11412
weighted avg	0.74	0.74	0.74	11412

test set precision recall f1-score support

0	0.97	0.72	0.82	531
1	0.17	0.71	0.27	42
accuracy			0.72	573
macro avg	0.57	0.71	0.55	573
weighted avg	0.91	0.72	0.78	573

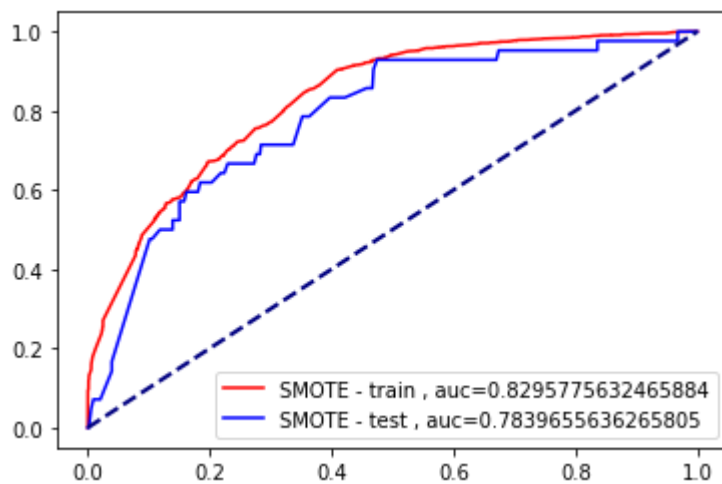
```

y_score_train = xgb_model_SMOTE.predict_proba(X_SMOTE)
fpr_train, tpr_train, _ = roc_curve(y_SMOTE, y_score_train[:,1])
auc_train = roc_auc_score(y_SMOTE, y_score_train[:,1])
plt.plot(fpr_train,tpr_train, color='red', label='SMOTE - train , auc='+str(auc_train))

y_score_test = xgb_model_SMOTE.predict_proba(X_test.to_numpy())
fpr_test, tpr_test, _ = roc_curve(y_test, y_score_test[:,1])
auc_test = roc_auc_score(y_test, y_score_test[:,1])
plt.plot(fpr_test,tpr_test, color='Blue', label='SMOTE - test , auc='+str(auc_test))

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.legend(loc=4)
plt.show()

```



✓ 0s completed at 17:36

● ×